CS 6220 Data Mining — Assignment 6

Regression

This assignment will require you to implement and interpret some of the regression concepts that were introduced in this module. Keep in mind that the main objective of this assignment is to highlight the insights that we can derive from applying these techniques—the coding aspect is secondary. Accordingly, you are welcome to consult any online documentation and/or code that has been posted to the course website, so long as all references and sources are properly cited. You are also encouraged to use code libraries, so long as you acknowledge any source code that was not written by you by mentioning the original author(s) directly in your source code (comment or header).
Using the Boston House-Prices dataset, you will be asked to construct a linear regression using different features of the dataset to predict the target. You will then be asked to evaluate the performance of the regression with visual outputs.

Objectives:

1. Apply linear regression to a dataset containing numerical features
2. Evaluate the performance of linear regression using R-squared metrics and Mean Squared Error

Submission:

Submit your ipynb file on the Assignment submission portal on Blackboard.

Grading Criteria:
Follow the instructions in the pdf, and complete each task. You will be graded on the application of the modules' topics, the completeness of your answers to the questions in the assignment notebook, and the clarity of your writing and code.

Assignment Description

The Data The given dataset contains information of Boston housing-prices. The dataset has 13 numeric/categorical predictive attributes. Median Value (attribute 14) is usually the target. You can know more about the dataset from here.
The dataset can be loaded using the sklearn datasets module. You cam get the feature names from the provided link.

The Idea: Linear Regression on Boston Hoise-Prices Dataset
Here, we want to predict the "Median value of owner-occupied homes" - that is the feature 14 of the data from other features. In this assignment, we will predict the target feature using all the other features. And then we will use single features to predict the target feature.

What to Do
First, load the dataset from sklearn dataset using the following code snippet.

```
from sklearn.datasets import load_boston
X, y = load_boston(return_X_y=True)
```

To generate a linear regression model, you may use the *linear_ model.LinearRegression()* function available via the scikit-learn library. To run the model on the Boston data, first divide the dataset into training and testing sets, then fit the model on the training set and predict with the fitted model on the testing set. scikit-learn provides several functions for dividing datasets in this manner, including *cross_ validation.KFold* and *cross_ validation.train_ test_ split*.
Several statistics can be generated from a linear model. Given a fitted linear model, the following code outputs the model coefficients (the parameter values for the fitted model), the residual sum of squares (the model error), and the explained variance (the degree to which the model explains the variation present in the data):

```
# The coefficients
print('Coefficients:', regr.coef )
# The mean squared error
print("Mean squared error:  %.2f" % np.mean((coly_pred  coly_test )  2))
# Explained variance score :  1 is perfect prediction
print('Variance score:  %.2f' % regr.score(colx_test , coly_test)
```

You can use these scores to measure the efficacy of a particular linear model.

First, use all the features (1-13) to fit the linear regression model for feature 14 using the training set. Report the r^2, mean squared error and variance score for the model on the test set.
Next, use each feature alone to fit a linear regression model on the same training set. Report the r^2, mean squared error and variance score for the model on the test set for each feature. Also show the visual plot.

What to Provide
Your output should contain the following:

- The evaluation of the linear regression model using all the features.
- The plot of the linear regression models generated on each feature wtih corresponding model statistics. The plot includes the training, test points and the linear regression line.
- Apply the same procedure (previous two steps) 10 times, and take the average values of the evaluation metrics (R^2 score, Mean Squared Error, Variance Score). Plot the average values to compare the model performances for each different criteria.

Given this output, respond to the following questions:

1. Based upon the linear models you generated, which feature appears to be most predictive for the target feature? Note that you can answer this question based upon the output provided for the linear models.
2. Suppose you need to select two features for a linear regression model to predict the target feature. Which two features are more likely to be selected? Why?
3. Compare the average evaluation metrics plot for different models.